

b)


```
=====
Training fCE for final 20 batches:
1: 2628.599011318614
2: 2831.6085264849107
3: 2639.7886243024077
4: 2609.741965900652
5: 2702.1394109395046
6: 2688.0205791754233
7: 2802.259330202891
8: 2629.6482559763995
9: 2710.908171366548
10: 2671.8496599937075
11: 2618.3492939401335
12: 2654.301574731994
13: 2608.384931469963
14: 2605.5610435320386
15: 2677.0171660785454
16: 2603.1886693343617
17: 2613.557061311925
18: 2648.385068122929
19: 2618.1782842580515
20: 2668.8688995582625
```

c)


```
=====
fPC:
83.49% Correct
```

Epoch count: 5; obtained through binary search


d)




w1:




w2:




$w_3$ :




$w_4$ :




w5:




w6:




$w_7$ :



$w_8$ :



w9:



w10:

Part 2:

My accuracy on the Kaggle competition was 0.77511. I used 15 epochs, and 100 batches